

Reproducibility & Environments

CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra, IIT Gandhinagar

The Reproducibility Crisis

"It works on my machine!"

Common scenarios:

- Research paper: "We achieved 95% accuracy"
- You: Run their code → 73% accuracy (or crash)
- Collaborator: "I can't install your dependencies"
- 6 months later: You can't run your own code

Root causes:

- Different Python versions
- Missing dependencies
- OS-specific issues
- Hardware differences
- Random seeds not set

What is Reproducibility?

Reproducibility: Ability to obtain same results using same code and data

Levels:

1. **Computational Reproducibility:** Same code + data = same results
2. **Statistical Reproducibility:** Control for randomness
3. **Scientific Reproducibility:** Independent verification

Why it matters:

- Scientific integrity
- Collaboration
- Debugging
- Deployment
- Future you will thank present you

Environment Management

The Problem:

Project A: Python 3.8, TensorFlow 2.4

Project B: Python 3.10, TensorFlow 2.12

System: Python 3.9, TensorFlow 2.8

Solution: Isolated environments

Tools:

- **venv/virtualenv**: Python virtual environments
- **conda**: Package and environment manager
- **Poetry**: Modern dependency management
- **Docker**: OS-level isolation

Python Virtual Environments (venv)

Create isolated Python environments

```
# Create virtual environment  
python -m venv myenv  
  
# Activate (Linux/Mac)  
source myenv/bin/activate  
  
# Activate (Windows)  
myenv\Scripts\activate  
  
# Install packages  
pip install numpy pandas scikit-learn  
  
# Save dependencies  
pip freeze > requirements.txt  
  
# Deactivate  
deactivate
```

requirements.txt

Standard way to specify dependencies

```
numpy=1.24.3
pandas=2.0.2
scikit-learn=1.2.2
matplotlib=3.7.1
torch=2.0.1
```

Install from requirements.txt:

```
pip install -r requirements.txt
```

Best practices:

- Pin versions (==) for reproducibility
- Use >= for flexibility
- Separate dev dependencies

Conda Environments

More powerful than venv, handles non-Python dependencies

```
# Create environment with specific Python version
conda create -n myenv python=3.10

# Activate
conda activate myenv

# Install packages
conda install numpy pandas scikit-learn

# Install from conda-forge
conda install -c conda-forge librosa

# Export environment
conda env export > environment.yml

# Create from YAML
conda env create -f environment.yml

# Deactivate
conda deactivate
```

environment.yml

Conda environment specification

```
name: ml-project
channels:
  - pytorch
  - conda-forge
  - defaults
dependencies:
  - python=3.10
  - numpy=1.24
  - pandas=2.0
  - pytorch=2.0
  - cudatoolkit=11.8
  - pip
  - pip:
    - transformers=4.30.0
    - wandb=0.15.0
```

Advantages:

Poetry: Modern Dependency Management

Handles dependencies, virtual environments, and packaging

```
# Install Poetry
curl -ssl https://install.python-poetry.org | python3 -

# Initialize project
poetry init

# Add dependencies
poetry add numpy pandas scikit-learn

# Add dev dependencies
poetry add --group dev pytest black

# Install all dependencies
poetry install

# Run command in environment
poetry run python train.py

# Build package
poetry build
```

pyproject.toml (Poetry)

```
[tool.poetry]
name = "ml-project"
version = "0.1.0"
description = "My ML project"
authors = ["Your Name <you@example.com>"]

[tool.poetry.dependencies]
python = "^3.10"
numpy = "^1.24"
pandas = "^2.0"
scikit-learn = "^1.2"
torch = "^2.0"

[tool.poetry.group.dev.dependencies]
pytest = "^7.3"
black = "^23.3"
mypy = "^1.3"

[build-system]
requires = ["poetry-core"]
build-backend = "poetry.core.masonry.api"
```

Setting Random Seeds

Ensure reproducible results

```
import random
import numpy as np
import torch

def set_seed(seed=42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)

    # Make CUDA deterministic
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

# Set at start of script
set_seed(42)

# Now your code is reproducible
```

What is Docker?

Containerization platform: Package code + dependencies + OS

Container vs VM:

Container	Virtual Machine
Lightweight (MB)	Heavy (GB)
Fast startup (seconds)	Slow startup (minutes)
Shares host kernel	Full OS
Less isolation	More isolation

Benefits:

- Works identically everywhere
- No "works on my machine" problems
- Easy to share and deploy

Docker Architecture

Key concepts:

- **Image:** Template for containers (like a class)
- **Container:** Running instance (like an object)
- **Dockerfile:** Instructions to build image
- **Registry:** Store images (Docker Hub, GitHub Container Registry)

Workflow:

```
Dockerfile → (build) → Image → (run) → Container
```

Dockerfile Basics

Create **Dockerfile**:

```
# Start from base image
FROM python:3.10-slim

# Set working directory
WORKDIR /app

# Copy requirements
COPY requirements.txt .

# Install dependencies
RUN pip install --no-cache-dir -r requirements.txt

# Copy project files
COPY ..

# Set environment variable
ENV PYTHONUNBUFFERED=1

# Run command
CMD ["python", "train.py"]
```

Building and Running Docker Images

Build image:

```
docker build -t my-ml-project:v1 .
```

Run container:

```
# Basic run  
docker run my-ml-project:v1  
  
# Interactive mode  
docker run -it my-ml-project:v1 /bin/bash  
  
# Mount volume (share files)  
docker run -v $(pwd)/data:/app/data my-ml-project:v1  
  
# Expose port (for APIs)  
docker run -p 8000:8000 my-ml-project:v1  
  
# Run with GPU  
docker run --gpus all my-ml-project:v1
```

Dockerfile for ML Projects

```
FROM pytorch/pytorch:2.0.1-cuda11.7-cudnn8-runtime

WORKDIR /app

# Install system dependencies
RUN apt-get update && apt-get install -y \
    git \
    wget \
&& rm -rf /var/lib/apt/lists/*

# Copy and install Python dependencies
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Copy project
COPY .

# Create directories
RUN mkdir -p /app/data /app/models /app/outputs

# Set Python path
ENV PYTHONPATH=/app

# Default command
CMD ["python", "train.py"]
```

Docker Compose

Manage multi-container applications

Create `docker-compose.yml`:

```
version: '3.8'

services:
  ml-training:
    build: .
    volumes:
      - ./data:/app/data
      - ./models:/app/models
    environment:
      - CUDA_VISIBLE_DEVICES=0
    deploy:
      resources:
        reservations:
          devices:
            - driver: nvidia
              count: 1
              capabilities: [gpu]

api:
  image: my-api:latest
  ports:
    - "8000:8000"
```

Using Docker Compose

```
# Start all services
docker-compose up

# Start in background
docker-compose up -d

# Stop services
docker-compose down

# View logs
docker-compose logs -f

# Run specific service
docker-compose run ml-training python train.py

# Rebuild images
docker-compose build
```

Multi-Stage Docker Builds

Reduce image size

```
# Stage 1: Build
FROM python:3.10 AS builder

WORKDIR /app
COPY requirements.txt .
RUN pip install --user --no-cache-dir -r requirements.txt

# Stage 2: Runtime
FROM python:3.10-slim

WORKDIR /app

# Copy installed packages from builder
COPY --from=builder /root/.local /root/.local

# Copy application
COPY ..

# Update PATH
ENV PATH=/root/.local/bin:$PATH

CMD ["python", "app.py"]
```

Docker Best Practices

1. Use specific base images

```
# Good  
FROM python:3.10-slim
```

```
# Avoid  
FROM python:latest
```

2. Minimize layers

```
# Good: Single RUN  
RUN apt-get update && apt-get install -y \  
    package1 package2 \  
    && rm -rf /var/lib/apt/lists/*
```

```
# Bad: Multiple RUNs  
RUN apt-get update  
RUN apt-get install -y package1  
RUN apt-get install -y package2
```

.dockerignore

Exclude files from Docker context

```
# .dockerignore
__pycache__
*.pyc
*.pyo
*.pyd
.Python
*.so
*.egg
*.egg-info
dist
build
.git
.gitignore
.vscode
.idea
*.ipynb_checkpoints
data/
models/*.pth
*.log
.env
```

Version Control for ML

Git + DVC (Data Version Control)

Problem: Git doesn't handle large files well

DVC: Version control for data and models

```
# Initialize DVC
dvc init

# Track large file
dvc add data/large_dataset.csv

# Commit
git add data/large_dataset.csv.dvc .gitignore
git commit -m "Add dataset"

# Push data to remote storage (S3, GCS, etc.)
dvc remote add -d myremote s3://mybucket/dvc-storage
dvc push

# Pull data
```

DVC Pipeline

Define reproducible ML pipelines

```
# dvc.yaml
stages:
  prepare:
    cmd: python prepare.py
    deps:
      - prepare.py
      - data/raw
    outs:
      - data/processed

  train:
    cmd: python train.py
    deps:
      - train.py
      - data/processed
    params:
      - train.learning_rate
      - train.epochs
    outs:
      - models/model.pkl
  metrics:
    - metrics.json:
        cache: false
```

Running DVC Pipelines

```
# Reproduce pipeline
dvc repro

# Visualize pipeline
dvc dag

# Compare experiments
dvc params diff

# Track metrics
dvc metrics show

# Create experiment
dvc exp run

# Compare experiments
dvc exp diff
```

MLflow for Experiment Tracking

Track experiments, models, and parameters

```
import mlflow
import mlflow.sklearn

# Start tracking
mlflow.set_tracking_uri("http://localhost:5000")
mlflow.set_experiment("my-experiment")

with mlflow.start_run():
    # Log parameters
    mlflow.log_param("learning_rate", 0.01)
    mlflow.log_param("epochs", 100)

    # Train model
    model = train_model()

    # Log metrics
    mlflow.log_metric("accuracy", 0.95)
    mlflow.log_metric("loss", 0.05)

    # Log model
    mlflow.sklearn.log_model(model, "model")

    # Log artifacts
    mlflow.log_artifact("plots/confusion_matrix.png")
```

MLflow UI

```
# Start MLflow server  
mlflow ui --host 0.0.0.0 --port 5000  
  
# Access at http://localhost:5000
```

Features:

- Compare experiment runs
- Visualize metrics
- Download models and artifacts
- Search and filter experiments
- Deploy models

Weights & Biases (wandb)

Modern experiment tracking and collaboration

```
import wandb

# Initialize
wandb.init(
    project="my-project",
    config={
        "learning_rate": 0.01,
        "epochs": 100,
        "batch_size": 32
    }
)

# Log metrics during training
for epoch in range(epochs):
    train_loss = train_epoch()
    val_loss = validate()

    wandb.log({
        "epoch": epoch,
        "train_loss": train_loss,
        "val_loss": val_loss
    })

# Log media
wandb.log({"confusion_matrix": wandb.plot.confusion_matrix( ... )})

# Finish
wandb.finish()
```

Configuration Management

Use config files instead of hardcoding

1. YAML/JSON configs:

```
# config.yaml
model:
    name: "resnet50"
    pretrained: true

training:
    batch_size: 32
    learning_rate: 0.001
    epochs: 100
    device: "cuda"

data:
    train_path: "data/train"
    val_path: "data/val"
    num_workers: 4
```

Loading Configs

```
import yaml
from dataclasses import dataclass

@dataclass
class TrainingConfig:
    batch_size: int
    learning_rate: float
    epochs: int
    device: str

def load_config(config_path):
    with open(config_path) as f:
        config = yaml.safe_load(f)

    training_config = TrainingConfig(**config['training'])
    return training_config

# Usage
config = load_config('config.yaml')
print(config.learning_rate)
```

Hydra for Configuration

Powerful configuration management

```
import hydra
from omegaconf import DictConfig

@hydra.main(config_path="configs", config_name="config", version_base=None)
def train(cfg: DictConfig):
    print(f"Learning rate: {cfg.training.learning_rate}")
    print(f"Batch size: {cfg.training.batch_size}")

    # Access nested config
    model = build_model(cfg.model)
    train_model(model, cfg.training)

if __name__ == "__main__":
    train()
```

Run with overrides:

```
python train.py training.learning_rate=0.01 training.epochs=50
```

Environment Variables

Store secrets and config

Create `.env` file:

```
API_KEY=your_secret_key  
DATABASE_URL=postgresql://localhost/db  
MODEL_PATH=/path/to/models  
DEBUG=True
```

Load in Python:

```
import os  
from dotenv import load_dotenv  
  
load_dotenv()  
  
api_key = os.getenv("API_KEY")  
db_url = os.getenv("DATABASE_URL")  
debug = os.getenv("DEBUG", "False") = "True"
```

Logging Best Practices

Structured logging for debugging

```
import logging

# Configure logging
logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
    handlers=[
        logging.FileHandler('training.log'),
        logging.StreamHandler()
    ]
)

logger = logging.getLogger(__name__)

# Use throughout code
logger.info("Starting training")
logger.debug(f"Batch size: {batch_size}")
logger.warning("Learning rate is very high")
logger.error("Failed to load checkpoint")

try:
    load_model()
except Exception as e:
    logger.exception("Error loading model")
```

Model Checkpointing

Save model state regularly

```
import torch

def save_checkpoint(model, optimizer, epoch, loss, path):
    checkpoint = {
        'epoch': epoch,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'loss': loss,
        'config': config
    }
    torch.save(checkpoint, path)
    logger.info(f"Checkpoint saved to {path}")

def load_checkpoint(model, optimizer, path):
    checkpoint = torch.load(path)
    model.load_state_dict(checkpoint['model_state_dict'])
    optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
    epoch = checkpoint['epoch']
    loss = checkpoint['loss']
    return model, optimizer, epoch, loss

# Save every N epochs
if epoch % 10 == 0:
    save_checkpoint(model, optimizer, epoch, loss,
                    f'checkpoints/model_epoch_{epoch}.pth')
```

Project Structure

Standard ML project layout:

```
ml-project/
├── data/
│   ├── raw/
│   ├── processed/
│   └── external/
├── models/
│   └── trained/
├── notebooks/
│   └── exploration.ipynb
└── src/
    ├── __init__.py
    ├── data/
    │   └── __init__.py
    │       └── dataset.py
    ├── models/
    │   └── __init__.py
    │       └── model.py
    └── utils/
        └── __init__.py
├── tests/
│   └── test_model.py
├── configs/
│   └── config.yaml
├── requirements.txt
├── Dockerfile
├── docker-compose.yml
├── .env.example
├── .gitignore
├── .dockerignore
├── README.md
└── setup.py
```

Reproducibility Checklist

Before sharing code:

- [] Pin all dependency versions
- [] Set random seeds
- [] Document Python version
- [] Include requirements.txt or environment.yml
- [] Provide Docker setup (Dockerfile)
- [] Document data sources
- [] Include sample data or download scripts
- [] Write clear README with setup instructions
- [] Test on clean environment
- [] Document hardware requirements (GPU, RAM)
- [] Include configuration files
- [] Version control with Git

README Template

```
# Project Name

## Setup

### Requirements
- Python 3.10+
- CUDA 11.7 (for GPU support)
- 16GB RAM minimum

### Installation

1. Clone repository
```
```
git clone https://github.com/user/project.git
cd project
```
```

2. Create virtual environment
```
```
python -m venv venv
source venv/bin/activate
```
```

3. Install dependencies
```
```
pip install -r requirements.txt
```
```

4. Download data
```
```
python scripts/download_data.py
```
```

## Usage

Train model:
```
```
python train.py --config configs/config.yaml
```
```

## Docker

```
```
docker build -t myproject .
docker run -v $(pwd)/data:/app/data myproject
```
```

```

Continuous Integration for ML

Automate testing and validation

`.github/workflows/train.yml :`

```
name: Train Model

on:
  push:
    branches: [ main ]

jobs:
  train:
    runs-on: ubuntu-latest

    steps:
      - uses: actions/checkout@v2

      - name: Set up Python
        uses: actions/setup-python@v2
        with:
          python-version: 3.10

      - name: Install dependencies
        run: |
          pip install -r requirements.txt

      - name: Run training
        run: |
          python train.py --epochs 5

      - name: Upload metrics
        uses: actions/upload-artifact@v2
        with:
          path: ./metrics
```

Model Cards

Document model details

```
# Model Card: Sentiment Classifier

## Model Details
- **Model type**: DistilBERT
- **Version**: 1.0
- **Date**: 2025-12-08

## Intended Use
- **Primary use**: Sentiment analysis of product reviews
- **Out-of-scope**: Political content, medical text

## Training Data
- **Dataset**: Amazon reviews (100k samples)
- **Splits**: 80/10/10 (train/val/test)
- **Preprocessing**: Lowercase, remove URLs

## Performance
- **Test Accuracy**: 92%
- **F1 Score**: 0.91

## Limitations
- Struggles with sarcasm
- Biased toward longer reviews

## Ethical Considerations
- May amplify existing biases in training data
```

What We've Learned

Environment Management:

- venv, conda, Poetry for Python environments
- requirements.txt, environment.yml, pyproject.toml

Containerization:

- Docker for reproducible environments
- Dockerfile, docker-compose
- Multi-stage builds

Version Control:

- Git for code
- DVC for data and models

Experiment Tracking:

- MLflow and Weights & Biases

Resources

Tools:

- Docker: <https://docs.docker.com/>
- Poetry: <https://python-poetry.org/>
- DVC: <https://dvc.org/>
- MLflow: <https://mlflow.org/>
- Weights & Biases: <https://wandb.ai/>

Guides:

- Reproducible ML: <https://www.nature.com/articles/s41592-021-01256-7>
- Docker for Data Science: <https://towardsdatascience.com/docker-for-data-science>
- DVC Tutorial: <https://dvc.org/doc/start>

Questions?

Next: Testing & CI/CD for AI

Lab: Dockerize ML project, setup experiment tracking