

# **Week 1 Lab: Web Scraping Workshop**

---

**CS 203: Software Tools and Techniques for AI**

**Duration:** 3 hours

**Prof. Nipun Batra & Teaching Assistants**

# Lab Overview

---

## Today's Goals

By the end of this lab, you will:

-  Master Chrome DevTools for finding data sources
-  Build scrapers with Requests + BeautifulSoup
-  Automate browsers with Playwright
-  Create a real data collection project
-  Understand ethical scraping practices

## Structure

- Part 1: DevTools Exploration (45 min)
- Part 2: Scraper Creation with BeautifulSoup (30 min)

# Setup Check (10 minutes)

---

## Verify Your Environment

```
# Check Python version (need 3.8+)
python --version

# Install required packages
pip install requests beautifulsoup4 playwright pandas matplotlib

# Install Playwright browsers
playwright install chromium

# Verify installations
python -c "import requests, bs4, playwright; print('All packages installed!')"
```

## Troubleshooting

- Permission errors: Use `pip install --user ...`

# Part 1: Chrome DevTools Mastery

---

## Exercise 1.1: Find the Hidden API (15 min)

Website: <https://www.goodreads.com> or <https://www.imdb.com>

Your Task:

1. Open the website in Chrome
2. Open DevTools (Cmd+Option+I / Ctrl+Shift+I)
3. Go to Network tab
4. Filter by XHR/Fetch
5. Search for a book/movie
6. Find the API request that returns search results
7. Copy the request as cURL

# Exercise 1.1: Solution Example

## IMDB Search Example

```
# What you might find in Network tab  
https://v3.sg.media-imdb.com/suggestion/x/avengers.json
```

```
# The response will be JSON like:
```

```
{  
  "d": [  
    {  
      "i": {"imageUrl": "..."},  
      "l": "Avengers: Endgame",  
      "q": "feature",  
      "rank": 43,  
      "s": "Robert Downey Jr., Chris Evans",  
      "y": 2019  
    },  
    ...  
  ]  
}
```

## Exercise 1.2: cURL to Python (15 min)

**Task:** Convert your cURL command to Python

### Step 1: Copy as cURL

Right-click request → Copy → Copy as cURL

### Step 2: Use curlconverter.com (or do manually)

- Visit <https://curlconverter.com>
- Paste your cURL command
- Select "Python Requests"
- Get Python code!

### Step 3: Test it

```
import requests
```

## Exercise 1.3: GitHub API Exploration (15 min)

**Task:** Use DevTools to find GitHub's API endpoints

1. Go to <https://github.com/microsoft/vscode>
2. Open DevTools → Network
3. Click on "Issues" tab
4. Find the API call for issues
5. Examine the response structure

**Questions to Answer:**

- What's the API endpoint?
- What parameters does it use?
- How is pagination handled?
- What authentication is needed?

# Part 1 Checkpoint

## What You've Learned

- How to find API endpoints using DevTools
- How to copy requests as cURL
- How to convert cURL to Python
- How to identify request parameters

## Quick Quiz (5 min)

1. What's the difference between XHR and Fetch?
2. Where do you find authentication tokens in DevTools?
3. Why is copying as cURL useful?

**Share your findings:** Each team shares one interesting API they found!

# Part 2: Static Scraping with BeautifulSoup

---

## Exercise 2.1: Quotes Scraper (20 min)

Website: <http://quotes.toscrape.com>

Task: Build a complete quotes scraper

```
import requests
from bs4 import BeautifulSoup

# TODO: Write code to scrape:
# 1. All quotes from the first page
# 2. Author of each quote
# 3. Tags for each quote
# 4. Save to a CSV file

# Starter code:
url = 'http://quotes.toscrape.com/'
response = requests.get(url)
```

# Exercise 2.1: Solution

```
import requests
from bs4 import BeautifulSoup
import csv

url = 'http://quotes.toscrape.com/'
response = requests.get(url)
soup = BeautifulSoup(response.text, 'html.parser')

quotes_data = []

# Find all quote containers
quotes = soup.find_all('div', class_='quote')

for quote in quotes:
    text = quote.find('span', class_='text').text
    author = quote.find('small', class_='author').text
    tags = [tag.text for tag in quote.find_all('a', class_='tag')]

    quotes_data.append({
        'quote': text,
        'author': author,
        'tags': ', '.join(tags)
    })

# Save to CSV
with open('quotes.csv', 'w', newline='', encoding='utf-8') as f:
    writer = csv.DictWriter(f, fieldnames=['quote', 'author', 'tags'])
    writer.writeheader()
    writer.writerows(quotes_data)

print(f"Saved {len(quotes_data)} quotes to quotes.csv")
```

## Exercise 2.2: Pagination (25 min)

**Task:** Extend your scraper to handle multiple pages

**Challenge:**

1. Scrape quotes from pages 1-10
2. Add a page number column
3. Include delay between requests
4. Handle the case when there are no more pages

**Hints:**

```
import time

base_url = 'http://quotes.toscrape.com'
page = 1

while True:
```

# Exercise 2.2: Solution

```
import requests
from bs4 import BeautifulSoup
import csv
import time

base_url = 'http://quotes.toscrape.com'
all_quotes = []
page = 1

while page <= 10: # Limit to 10 pages
    url = f'{base_url}/page/{page}/'
    print(f"Scraping page {page}...")

    response = requests.get(url)
    soup = BeautifulSoup(response.text, 'html.parser')

    quotes = soup.find_all('div', class_='quote')

    if not quotes: # No more quotes
        print("No more pages!")
        break

    for quote in quotes:
        text = quote.find('span', class_='text').text
        author = quote.find('small', class_='author').text
        tags = [tag.text for tag in quote.find_all('a', class_='tag')]

        all_quotes.append({
            'page': page,
            'quote': text,
            'author': author,
            'tags': ', '.join(tags)
        })

    page += 1
    time.sleep(1) # Polite delay

# Save to CSV
with open('all_quotes.csv', 'w', newline='', encoding='utf-8') as f:
    writer = csv.DictWriter(f, fieldnames=['page', 'quote', 'author', 'tags'])
    writer.writeheader()
    writer.writerows(all_quotes)

print(f"Total quotes scraped: {len(all_quotes)})")
```

## Exercise 2.3: News Aggregator (45 min)

**Task:** Build a news headline aggregator

**Website Options** (choose one):

- <https://news.ycombinator.com> (Hacker News)
- <https://www.reddit.com/r/technology> (Reddit)
- <https://www.bbc.com/news> (BBC News)

**Requirements:**

1. Scrape headlines and links
2. Extract publish date/time if available
3. Get vote count or popularity metric
4. Save to CSV
5. Create a simple visualization with pandas

# Exercise 2.3: Hacker News Solution

```
import requests
from bs4 import BeautifulSoup
import pandas as pd
import matplotlib.pyplot as plt

url = 'https://news.ycombinator.com/'
response = requests.get(url)
soup = BeautifulSoup(response.text, 'html.parser')

stories = []

# Find all story rows
story_links = soup.find_all('span', class_='titleline')
subtexts = soup.find_all('td', class_='subtext')

for link_span, subtext in zip(story_links, subtexts):
    link = link_span.find('a')
    title = link.text
    url = link['href']

    # Get points
    score = subtext.find('span', class_='score')
    points = int(score.text.split()[0]) if score else 0

    # Get author and time
    author = subtext.find('a', class_='hnuser')
    author_name = author.text if author else 'Unknown'

    stories.append({
        'title': title,
        'url': url,
        'points': points,
        'author': author_name
    })

# Create DataFrame
df = pd.DataFrame(stories)

# Save to CSV
df.to_csv('hackernews.csv', index=False)

# Visualize top stories
df.nlargest(10, 'points').plot(
    x='title', y='points', kind='barh',
    figsize=(10, 6), title='Top 10 HN Stories'
)
plt.tight_layout()
plt.savefig('top_stories.png')

print(f"Scraped {len(stories)} stories")
print(f"\nTop 5 stories:")
print(df.nlargest(5, 'points')[['title', 'points']])
```

# Data Cleaning & Analysis (10 min)

Task: Clean your scraped data with pandas

```
import pandas as pd

df = pd.read_csv('hackernews.csv')

# Clean data
df['title'] = df['title'].str.strip()
df['points'] = df['points'].fillna(0).astype(int)

# Analysis
print("Statistics:")
print(df['points'].describe())

print("\nMost popular domains:")
df['domain'] = df['url'].str.extract(r'https?:\/\/([^\s]+)')
print(df['domain'].value_counts().head())

print("\nMost prolific authors:")
print(df['author'].value_counts().head())

# Word frequency in titles
from collections import Counter
all_words = ' '.join(df['title']).lower().split()
common_words = Counter(all_words).most_common(20)
```

## Part 2 Checkpoint

### What You've Built

- Quotes scraper with pagination
- News aggregator with data analysis
- CSV export and pandas integration
- Basic visualization

### Discussion (10 min)

#### Questions:

1. What was the hardest part of scraping?
2. Did you encounter any errors? How did you fix them?
3. What patterns did you notice in HTML structure?

**Challenge:** Who scraped the most data? Share your stats!

# Part 3: Playwright for Dynamic Content

---

## Exercise 3.1: First Playwright Script (15 min)

Task: Scrape Google search results

```
from playwright.sync_api import sync_playwright

with sync_playwright() as p:
    # Launch browser (headless=False to see what's happening)
    browser = p.chromium.launch(headless=False)
    page = browser.new_page()

    # Search on Google
    page.goto('https://www.google.com')

    # TODO:
    # 1. Find search box and type "web scraping python"
    # 2. Click search button
    # 3. Wait for results
```

# Exercise 3.1: Solution

```
from playwright.sync_api import sync_playwright
import time

with sync_playwright() as p:
    browser = p.chromium.launch(headless=False)
    page = browser.new_page()

    # Go to Google
    page.goto('https://www.google.com')

    # Search
    page.fill('textarea[name="q"]', 'web scraping python')
    page.press('textarea[name="q"]', 'Enter')

    # Wait for results
    page.wait_for_selector('h3')

    # Extract results
    results = page.query_selector_all('h3')

    print("Top 10 results:")
    for i, result in enumerate(results[:10], 1):
        title = result.inner_text()
        # Get parent link
        link_elem = result.query_selector('xpath=ancestor::a')
        url = link_elem.get_attribute('href') if link_elem else 'No URL'

        print(f'{i}. {title}')
        print(f'  {url}\n')

    time.sleep(2) # Pause to see results
    browser.close()
```

## Exercise 3.2: Infinite Scroll (15 min)

**Task:** Scrape content from an infinite scroll page

**Website:** <https://www.reddit.com> or similar

```
from playwright.sync_api import sync_playwright
import time

with sync_playwright() as p:
    browser = p.chromium.launch(headless=False)
    page = browser.new_page()

    page.goto('https://www.reddit.com/r/programming/')

    # TODO:
    # 1. Scroll down 5 times to load more posts
    # 2. Wait between scrolls for content to load
    # 3. Extract all post titles
    # 4. Save to CSV

    # Hint: use page.evaluate() to scroll
```

# Exercise 3.2: Solution

```
from playwright.sync_api import sync_playwright
import time
import csv

with sync_playwright() as p:
    browser = p.chromium.launch(headless=False)
    page = browser.new_page()

    page.goto('https://www.reddit.com/r/programming/')

    # Wait for initial content
    page.wait_for_selector('h3')

    # Scroll 5 times
    for i in range(5):
        print(f"Scroll {i+1}/5...")

        # Get current height
        previous_height = page.evaluate('document.body.scrollHeight')

        # Scroll to bottom
        page.evaluate('window.scrollTo(0, document.body.scrollHeight)')

        # Wait for new content
        time.sleep(2)

        # Wait for height to change (new content loaded)
        new_height = page.evaluate('document.body.scrollHeight')
        if new_height == previous_height:
            print('No more content to load')
            break

    # Extract all posts
    posts = page.query_selector_all('h3')

    results = []
    for post in posts:
        title = post.inner_text()
        results.append({'title': title})

    # Save to CSV
    with open('reddit_posts.csv', 'w', newline='', encoding='utf-8') as f:
        writer = csv.DictWriter(f, fieldnames=['title'])
        writer.writeheader()
        writer.writerows(results)

    print(f"Scraped {len(results)} posts")
    browser.close()
```

## Exercise 3.3: Login & Authentication (15 min)

Task: Automate login to a website

Practice Site: <http://quotes.toscrape.com/login>

```
from playwright.sync_api import sync_playwright

with sync_playwright() as p:
    browser = p.chromium.launch(headless=False)
    page = browser.new_page()

    # Go to login page
    page.goto('http://quotes.toscrape.com/login')

    # TODO:
    # 1. Fill in username (any username works)
    # 2. Fill in password (any password works)
    # 3. Click login button
    # 4. Verify you're logged in
    # 5. Navigate to a protected page
    # 6. See if you can still see the protected content
```

# Exercise 3.3: Solution

```
from playwright.sync_api import sync_playwright
import time

with sync_playwright() as p:
    browser = p.chromium.launch(headless=False)
    page = browser.new_page()

    # Login
    page.goto('http://quotes.toscrape.com/login')

    page.fill('input[name="username"]', 'admin')
    page.fill('input[name="password"]', 'admin')
    page.click('input[type="submit"]')

    # Wait for redirect
    page.wait_for_url('**/quotes.toscrape.com/')

    # Verify login
    if page.query_selector('a[href="/logout"]'):
        print('/ Successfully logged in!')
    else:
        print("x Login failed")

    # Now scrape protected content
    response = page.goto('http://quotes.toscrape.com/')

    # Extract quotes (same as before, but now we're authenticated)
    quotes = page.query_selector_all('div.quote')

    for quote in quotes[:5]:
        text = quote.query_selector('span.text').inner_text()
        author = quote.query_selector('small.author').inner_text()
        print(f'{text}\n - {author}\n')

    time.sleep(2)
    browser.close()
```

## Part 3 Checkpoint

### Playwright Skills Acquired

- Browser automation basics
- Form filling and clicking
- Infinite scroll handling
- Login automation
- Dynamic content extraction

### When to Use Playwright

Use Playwright when:

- Content loads via JavaScript
- Need to interact with page
- Infinite scroll or lazy loading

# Part 4: Mini Project (30 minutes)

---

## Your Own Scraping Project!

**Task:** Choose a website and build a complete scraper

### Requirements

1. **Choose** a website you're interested in
2. **Decide** what data to collect
3. **Build** a scraper (Requests+BS4 or Playwright)
4. **Collect** at least 50 data points
5. **Clean** and analyze with pandas
6. **Visualize** one insight
7. **Present** to the class (2 min each)

# Project Template

```
"""
Project: [Your Project Name]
Goal: [What you're scraping]
Website: [URL]
"""

import requests
from bs4 import BeautifulSoup
import pandas as pd
import matplotlib.pyplot as plt

# Configuration
BASE_URL = 'https://...'
OUTPUT_FILE = 'data.csv'

# 1. Scraping function
def scrape_data():
    # Your scraping code here
    data = []
    # ...
    return data

# 2. Data cleaning
def clean_data(raw_data):
    df = pd.DataFrame(raw_data)
    # Clean data...
    return df

# 3. Analysis & Visualization
def analyze(df):
    print(df.describe())
    # Create plot...

# 4. Main execution
if __name__ == '__main__':
    raw_data = scrape_data()
    df = clean_data(raw_data)
    df.to_csv(OUTPUT_FILE, index=False)
    analyze(df)
```

# Project Rubric

## Grading Criteria (for reference)

Criteria	Points
Data collection works	30%
Code quality & comments	20%
Data cleaning	15%
Analysis/Visualization	15%
Follows ethical practices	10%
Presentation	10%

## Ethical Checklist

-  Checked robots.txt

# Case Study 1: Price Tracking

## Real-World Application

**Scenario:** Track laptop prices across e-commerce sites

```
import requests
from bs4 import BeautifulSoup
import pandas as pd
from datetime import datetime

def scrape_amazon_price(product_url):
    headers = {'User-Agent': 'Mozilla/5.0...'}
    response = requests.get(product_url, headers=headers)
    soup = BeautifulSoup(response.text, 'html.parser')

    # Find price element
    price = soup.find('span', class_='a-price-whole')
    if price:
        return float(price.text.replace(',', ''))
    return None

# Track multiple products
products = {
    'Laptop A': 'https://amazon.in/...',
    'Laptop B': 'https://flipkart.com/...'
}

# Scrape daily and save to CSV
price_data = []
for name, url in products.items():
    price = scrape_amazon_price(url)
    price_data.append({
        'date': datetime.now(),
        'product': name,
        'price': price
    })
```

# Case Study 2: Research Paper Monitoring

## Academic Use Case

Scenario: Track new papers on arXiv by topic

```
import requests
from bs4 import BeautifulSoup
import re

def scrape_arxiv_recent(category='cs.AI', max_results=10):
    """Scrape recent papers from arXiv"""

    url = f'https://arxiv.org/list/{category}/recent'
    response = requests.get(url)
    soup = BeautifulSoup(response.text, 'html.parser')

    papers = []
    paper_divs = soup.find_all('dt')
    desc_divs = soup.find_all('dd')

    for paper, desc in zip(paper_divs[:max_results], desc_divs[:max_results]):
        # Extract arXiv ID
        arxiv_id = paper.find('a', title= Abstract')
        if not arxiv_id:
            continue

        arxiv_id = arxiv_id.text.strip()

        # Extract title and authors
        title = desc.find('div', class_='list-title').text
        title = re.sub(r'^Title:\s*', '', title).strip()

        authors = desc.find('div', class_='list-authors').text
        authors = re.sub(r'^Authors:\s*', '', authors).strip()

        papers.append({
            'arxiv_id': arxiv_id,
            'title': title,
            'authors': authors,
            'url': f'https://arxiv.org/abs/{arxiv_id}'
        })
```

# Case Study 3: Social Media Monitoring

## Marketing Use Case

### Scenario: Track brand mentions on Reddit

```
from playwright.sync_api import sync_playwright
import time

def monitor_reddit_mentions(keyword, subreddit='all', limit=20):
    """Monitor Reddit for keyword mentions"""

    with sync_playwright() as p:
        browser = p.chromium.launch(headless=True)
        page = browser.new_page()

        # Search Reddit
        search_url = f"https://www.reddit.com/search/?q={keyword}&sort=new"
        page.goto(search_url)

        # Wait for results
        page.wait_for_selector('h3', timeout=10000)

        # Scroll to load more
        for _ in range(3):
            page.evaluate("window.scrollTo(0, document.body.scrollHeight)")
            time.sleep(2)

        # Extract posts
        posts = page.query_selector_all('div[data-testid="post-container"]')

        mentions = []
        for post in posts[:limit]:
            try:
                title_elem = post.query_selector('h3')
                if not title_elem:
                    continue

                title = title_elem.inner_text()

                # Get metadata
                author = post.query_selector('tdata-testid="post_author_link"')
                author = author.inner_text() if author else 'Unknown'

                # Get score
                score = post.query_selector('[id^="vote-arrows"]')
                score_text = score.inner_text() if score else '0'

                mentions.append({
                    'title': title,
                    'author': author,
                    'score': score_text
                })
            except Exception as e:
                print(f"Error processing post: {e}")

    browser.close()
```

# Debugging Common Issues

## Problem 1: Empty Results

```
# Issue: soup.find_all() returns empty list
soup.find_all('div', class_='item') # []

# Debug steps:
print(soup.prettify()) # See actual HTML
print(len(response.text)) # Check if page loaded

# Common causes:
# 1. Wrong selector
# 2. JavaScript-rendered content (use Playwright)
# 3. Need to login first
# 4. Blocked by anti-scraping (add headers)
```

## Problem 2: 403 Forbidden

```
# Issue: requests.get() returns 403
```

# Debugging Common Issues (Continued)

## Problem 3: Element Not Found

```
# Issue: AttributeError: 'NoneType' object has no attribute 'text'

# Bad code:
title = soup.find('h1').text # Crashes if h1 not found

# Good code:
title_elem = soup.find('h1')
if title_elem:
    title = title_elem.text
else:
    title = 'No title found'

# Or use get_text with default:
title = soup.find('h1').get_text(default='No title') if soup.find('h1') else 'No title'

# Better: Use try-except
try:
    title = soup.find('h1').text
except AttributeError:
    title = 'No title'
```

# Best Practices Checklist

## Before You Scrape

- [ ] Check if API exists (easier than scraping!)
- [ ] Read `robots.txt` (`site.com/robots.txt`)
- [ ] Review Terms of Service
- [ ] Check data usage licenses
- [ ] Plan your rate limiting strategy

## While Scraping

- [ ] Use descriptive User-Agent
- [ ] Add delays between requests (1-2s minimum)
- [ ] Handle errors gracefully
- [ ] Log your activities

# Presentation Guidelines

## What to Present (2 minutes)

1. What did you scrape? (10 sec)

- Website and data type

2. How did you scrape it? (30 sec)

- Requests or Playwright?
- Main challenges overcome

3. What did you find? (60 sec)

- Show your data (table or visualization)
- One interesting insight

4. Code snippet (20 sec)

# Sample Presentation Structure

## # Project: IITGN Mess Menu Scraper

### ## What

Scraped weekly mess menu from IITGN website

### ## How

- Used Requests + BeautifulSoup
- Challenge: Menu in table format with merged cells
- Solution: Used pandas.read\_html()

### ## Findings

- 73 unique dishes over 4 weeks
- Most common: "Dal Fry" (appears 12 times)
- Visualization: Word cloud of dishes

### ## Code Highlight

```
```python
# pandas.read_html() saved the day!
tables = pd.read_html(url)
menu_df = tables[0] # First table is the menu
```

# Resources for Continued Learning

## Documentation

- **Requests:** <https://docs.python-requests.org/>
- **BeautifulSoup:** <https://www.crummy.com/software/BeautifulSoup/>
- **Playwright:** <https://playwright.dev/python/>
- **pandas:** <https://pandas.pydata.org/docs/>

## Practice Sites

- <http://quotes.toscrape.com> - Practice scraping
- <http://books.toscrape.com> - E-commerce practice
- <https://scrapethissite.com> - Various challenges
- <https://webscraper.io/test-sites> - Test different patterns

# Lab Submission

## What to Submit

### 1. Code (Python files or Jupyter notebook)

- Well-commented
- Includes requirements.txt

### 2. Data (CSV file)

- Cleaned and formatted
- At least 50 data points

### 3. Report (README.md or PDF)

- What you scraped
- How you did it

# Homework for Next Week

## Assignments

### 1. Extend your scraper

- Collect 200+ data points
- Add error handling
- Create 3 visualizations

### 2. Read

- Requests documentation (focus on sessions)
- BeautifulSoup CSS selectors guide
- Web scraping ethics article (link on course site)

### 3. Install for next week

# Great Work Today! 🎉

---

## You've Learned:

- Chrome DevTools for finding APIs
- Web scraping with Requests + BeautifulSoup
- Browser automation with Playwright
- Data collection, cleaning, and analysis

## Next Week:

Data Validation, Labeling, and Quality Control

Questions? Office hours tomorrow 3-5 PM

# Quick Feedback (5 min)

## Anonymous Poll

1. **Pace:** Too fast / Just right / Too slow
2. **Difficulty:** Too easy / Just right / Too hard
3. **Most useful:** DevTools / Requests / Playwright
4. **What to improve:** [Open feedback]

**Scan QR code or visit:**

[feedback link]

**Thank you! See you next week!**